

# Kaggle Playground

## Problem Statement / Real World Implementations

### 1. Importing Libraries

```
In [1]: # Core Data Science Libraries
import numpy as np
import pandas as pd
import warnings

# Visualization Libraries
import plotly.express as px
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.graph_objects as go
from plotly.subplots import make_subplots

# Scikit-Learn for Preprocessing and Modeling
from sklearn.model_selection import KFold, train_test_split
from sklearn.preprocessing import OrdinalEncoder, StandardScaler
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Machine Learning Models
from xgboost import XGBRegressor
from lightgbm import LGBMRegressor

# Hyperparameter Tuning
import optuna

# Notebook settings
warnings.filterwarnings('ignore')
pd.set_option('display.max_columns', None)
```

### 2. Loading Dataset

```
In [2]: # Define file paths
TRAIN_PATH = "/kaggle/input/playground-series-s5e11/train.csv"
TEST_PATH = "/kaggle/input/playground-series-s5e11/test.csv"
SUBMISSION_PATH = "/kaggle/input/playground-series-s5e11/sample_submission.csv"

# Load the datasets into pandas DataFrames
train_df = pd.read_csv(TRAIN_PATH)
test_df = pd.read_csv(TEST_PATH)
submission_df = pd.read_csv(SUBMISSION_PATH)
```

### Refining the Original dataset

```
In [3]: print("Train shape:", train_df.shape)
print("Test shape:", test_df.shape)
```

```
Train shape: (593994, 13)
Test shape: (254569, 12)
```

```
In [4]: df=train_df
df.head(5)
```

Out[4]:

	<b>id</b>	<b>annual_income</b>	<b>debt_to_income_ratio</b>	<b>credit_score</b>	<b>loan_amount</b>	<b>interest_rate</b>	<b>gend</b>
<b>0</b>	0	29367.99	0.084	736	2528.42	13.67	Fema
<b>1</b>	1	22108.02	0.166	636	4593.10	12.92	Ma
<b>2</b>	2	49566.20	0.097	694	17005.15	9.76	Ma
<b>3</b>	3	46858.25	0.065	533	4682.48	16.10	Fema
<b>4</b>	4	25496.70	0.053	665	12184.43	10.21	Ma

In [5]:

```
print(df["gender"].unique())
print(df["marital_status"].unique())
print(df["education_level"].unique())
print(df["employment_status"].unique())
print(df["loan_purpose"].unique())
print(df["grade_subgrade"].unique())
```

```
['Female' 'Male' 'Other']
['Single' 'Married' 'Divorced' 'Widowed']
['High School' "Master's" "Bachelor's" 'PhD' 'Other']
['Self-employed' 'Employed' 'Unemployed' 'Retired' 'Student']
['Other' 'Debt consolidation' 'Home' 'Education' 'Vacation' 'Car'
 'Medical' 'Business']
['C3' 'D3' 'C5' 'F1' 'D1' 'D5' 'C2' 'C1' 'F5' 'D4' 'C4' 'D2' 'E5' 'B1'
 'B2' 'F4' 'A4' 'E1' 'F2' 'B4' 'E4' 'B3' 'E3' 'B5' 'E2' 'F3' 'A5' 'A3'
 'A1' 'A2']
```

In [6]:

```
df.isna().sum()
```

Out[6]:

```
id                0
annual_income     0
debt_to_income_ratio 0
credit_score      0
loan_amount        0
interest_rate      0
gender              0
marital_status      0
education_level      0
employment_status      0
loan_purpose        0
grade_subgrade       0
loan_paid_back       0
dtype: int64
```

In [7]:

```
df.head(10)
```

Out[7]:

	<b>id</b>	<b>annual_income</b>	<b>debt_to_income_ratio</b>	<b>credit_score</b>	<b>loan_amount</b>	<b>interest_rate</b>	<b>gend</b>
<b>0</b>	0	29367.99	0.084	736	2528.42	13.67	Fema
<b>1</b>	1	22108.02	0.166	636	4593.10	12.92	Ma
<b>2</b>	2	49566.20	0.097	694	17005.15	9.76	Ma
<b>3</b>	3	46858.25	0.065	533	4682.48	16.10	Fema
<b>4</b>	4	25496.70	0.053	665	12184.43	10.21	Ma
<b>5</b>	5	44940.30	0.058	653	12159.92	12.24	Ma
<b>6</b>	6	61574.16	0.042	696	16907.71	13.52	Oth
<b>7</b>	7	45953.31	0.100	654	10111.62	12.82	Fema
<b>8</b>	8	30592.29	0.132	713	7522.36	9.48	Ma
<b>9</b>	9	17342.45	0.121	548	9653.48	16.04	Fema

## 4. EDA

In [8]:

```
# Select only numeric columns for correlation matrix
numerical_cols = train_df.select_dtypes(include=np.number).columns.tolist()
numerical_cols.remove('id')
numerical_cols.remove('loan_paid_back')

numeric_df = train_df[numerical_cols + ['loan_paid_back']]
corr_matrix = numeric_df.corr()

# Create the interactive heatmap
fig = go.Figure(data=go.Heatmap(
    z=corr_matrix.values,
    x=corr_matrix.columns,
    y=corr_matrix.columns,
    colorscale='RdBu_r',
    zmin=-1, zmax=1,
    text=corr_matrix.round(2).values,
    texttemplate="%{text}",
    hoverongaps=False))

fig.update_layout(
    title='Correlation Heatmap of Numerical Features',
    width=800, height=800
)
fig.show()
```

### 3. Normalization of data

```
In [9]: # Save the test IDs for your submission file
test_ids = test_df['id']

# Drop the 'id' column from both, as it's not a feature
train_df = train_df.drop('id', axis=1)
test_df = test_df.drop('id', axis=1)
```

```
In [10]: def create_financial_features(df):
    """
```

```

Creates new financial features from the existing columns.
"""

# Create monthly_income
df['monthly_income'] = df['annual_income'] / 12

# Create total_monthly_debt
df['total_monthly_debt'] = df['debt_to_income_ratio'] * df['monthly_income']

# Create available_income (disposable income)
df['available_income'] = df['monthly_income'] - df['total_monthly_debt']

# Create loan_to_income_ratio
df['loan_to_income_ratio'] = df['loan_amount'] / df['annual_income']

# Create loan_to_available_income
# We will replace inf/-inf with NaN later
df['loan_to_available_income'] = df['loan_amount'] / df['available_income']

return df

# Assuming 'train_df' and 'test_df' are already loaded

print("Applying features to train_df...")
train_df = create_financial_features(train_df)

print("Applying features to test_df...")
test_df = create_financial_features(test_df)

```

Applying features to train\_df...

Applying features to test\_df...

```

In [11]: import pandas as pd
import numpy as np

def process_risk_features(df):
    """
    Applies ordinal mapping and binning to risk features.
    (Corrected version)
    """

    # 1. Map grade_subgrade
    # This map is correct and complete.
    grades = ['A', 'B', 'C', 'D', 'E', 'F', 'G']
    subgrades = ['1', '2', '3', '4', '5']
    grade_map = {}
    i = 0
    for g in grades:
        for s in subgrades:
            grade_map[g + s] = i
            i += 1

    df['grade_subgrade_ordinal'] = df['grade_subgrade'].map(grade_map)

    # 2. Map education_Level (FIXED)
    # This map now includes 'PhD' and 'Other'
    # We place them in a logical ordinal scale.
    education_map = {
        'Other': 0,           # Treat 'Other' as the baseline
        'High School': 1,
        'Bachelor\'s': 2,
    }

```

```

        'Master\'s': 3,
        'PhD': 4           # 'PhD' is the highest
    }
df['education_level_ordinal'] = df['education_level'].map(education_map)

# 3. Bin credit_score
# This logic is correct.
score_bins = [300, 579, 669, 739, 799, 850]
score_labels = ['Poor', 'Fair', 'Good', 'Very Good', 'Excellent']

df['credit_score_bin'] = pd.cut(df['credit_score'],
                                bins=score_bins,
                                labels=score_labels,
                                include_lowest=True)

# 4. Drop original columns
cols_to_drop = ['grade_subgrade', 'education_level', 'credit_score']
df = df.drop(columns=cols_to_drop)

return df

```

In [12]: cols\_to\_drop = ['annual\_income', 'debt\_to\_income\_ratio']

```

# --- Clean Training Data ---
train_df = train_df.drop(columns=cols_to_drop)
train_df.replace([np.inf, -np.inf], np.nan, inplace=True)

# --- Clean Test Data ---
test_df = test_df.drop(columns=cols_to_drop)
test_df.replace([np.inf, -np.inf], np.nan, inplace=True)

print("Feature creation complete. Redundant columns dropped.")
print("\n--- train_df columns ---")
print(list(train_df.columns))

print("\n--- test_df columns ---")
print(list(test_df.columns))

```

Feature creation complete. Redundant columns dropped.

```

--- train_df columns ---
['credit_score', 'loan_amount', 'interest_rate', 'gender', 'marital_status',
'education_level', 'employment_status', 'loan_purpose', 'grade_subgrade', 'loan_paid',
'monthly_income', 'total_monthly_debt', 'available_income', 'loan_to_income_ratio',
'loan_to_available_income']

--- test_df columns ---
['credit_score', 'loan_amount', 'interest_rate', 'gender', 'marital_status',
'education_level', 'employment_status', 'loan_purpose', 'grade_subgrade', 'monthly_ir',
'total_monthly_debt', 'available_income', 'loan_to_income_ratio',
'loan_to_available_income']

```

In [13]: # --- Apply to train\_df ---
print("Processing train\_df...")
train\_df = process\_risk\_features(train\_df)

# --- Apply to test\_df ---
print("Processing test\_df...")
test\_df = process\_risk\_features(test\_df)

```

print("\nFeature creation complete. Original columns dropped.")

# --- Show New Columns ---
print("\n--- train_df columns ---")
print(list(train_df.columns))

print("\n--- test_df columns ---")
print(list(test_df.columns))

# --- Check the new binned values ---
print("\n--- Value counts for new 'credit_score_bin' (from train_df) ---")
print(train_df['credit_score_bin'].value_counts())

```

Processing train\_df...

Processing test\_df...

Feature creation complete. Original columns dropped.

```

--- train_df columns ---
['loan_amount', 'interest_rate', 'gender', 'marital_status', 'employment_status',
'loan_purpose', 'loan_paid_back', 'monthly_income', 'total_monthly_debt',
'available_income', 'loan_to_income_ratio', 'loan_to_available_income',
'grade_subgrade_ordinal', 'education_level_ordinal', 'credit_score_bin']

--- test_df columns ---
['loan_amount', 'interest_rate', 'gender', 'marital_status', 'employment_status',
'loan_purpose', 'monthly_income', 'total_monthly_debt', 'available_income',
'loan_to_income_ratio', 'loan_to_available_income', 'grade_subgrade_ordinal',
'education_level_ordinal', 'credit_score_bin']

--- Value counts for new 'credit_score_bin' (from train_df) ---
credit_score_bin
Good          275585
Fair           209954
Very Good      71337
Poor            27266
Excellent       9852
Name: count, dtype: int64

```

In [14]: `from sklearn.preprocessing import OneHotEncoder`

```

categorical_cols = [
    'gender',
    'marital_status',
    'employment_status',
    'loan_purpose',
    'credit_score_bin'
]

ohe = OneHotEncoder(handle_unknown='ignore', sparse_output=False, dtype='int')

print("Fitting One-Hot Encoder on train_df...")
# Note: We must convert .astype(str) because 'credit_score_bin' is a category
ohe.fit(train_df[categorical_cols].astype(str))

# 3. TRANSFORM both train_df and test_df
print("Transforming train_df...")
# --- Process train_df ---
# Get the new encoded columns

```

```

train_encoded = ohe.transform(train_df[categorical_cols].astype(str))
# Get the new column names
encoded_feature_names = ohe.get_feature_names_out(categorical_cols)
# Create a new DataFrame with the encoded features
train_encoded_df = pd.DataFrame(train_encoded, columns=encoded_feature_names, index=index)
# Drop the original text columns
train_processed = train_df.drop(columns=categorical_cols)
# Add the new numeric columns
train_processed = pd.concat([train_processed, train_encoded_df], axis=1)

print("Transforming test_df...")
# --- Process test_df ---
# Use the SAME fitted encoder to transform test_df
test_encoded = ohe.transform(test_df[categorical_cols].astype(str))
# Create a new DataFrame with the encoded features
test_encoded_df = pd.DataFrame(test_encoded, columns=encoded_feature_names, index=index)
# Drop the original text columns
test_processed = test_df.drop(columns=categorical_cols)
# Add the new numeric columns
test_processed = pd.concat([test_processed, test_encoded_df], axis=1)

print("\n--- Encoding Complete ---")
print("train_processed is ready for training.")
print("test_processed is ready for prediction.")

```

Fitting One-Hot Encoder on train\_df...

Transforming train\_df...

Transforming test\_df...

--- Encoding Complete ---

train\_processed is ready for training.

test\_processed is ready for prediction.

In [15]: `train_df.head(5)`

	loan_amount	interest_rate	gender	marital_status	employment_status	loan_purpose	lc
<b>0</b>	2528.42	13.67	Female	Single	Self-employed		Other
<b>1</b>	4593.10	12.92	Male	Married	Employed		Debt consolidation
<b>2</b>	17005.15	9.76	Male	Single	Employed		Debt consolidation
<b>3</b>	4682.48	16.10	Female	Single	Employed		Debt consolidation
<b>4</b>	12184.43	10.21	Male	Married	Employed		Other

In [16]: `# Exclude target column if present`  
`features = train_processed.drop(columns=['loan_paid_back'], errors='ignore')`

`# 1. Check summary statistics`  
`print("Summary Statistics:\n")`  
`display(features.describe())`

`# 2. Check for large differences in scale`

```

range_df = features.max() - features.min()
print("\nFeature Ranges:\n")
print(range_df.sort_values(ascending=False))

# 3. Visualize distribution of feature scales
plt.figure(figsize=(10, 6))
sns.boxplot(data=features, orient='h', fliersize=1)
plt.title("Feature Value Distributions (Check for Scale Differences)")
plt.xscale('log')
plt.show()

# 4. Correlation check
corr_matrix = features.corr()
high_range_features = range_df[range_df > range_df.mean()].index.tolist()
print(f"\nFeatures with significantly higher ranges: {high_range_features}")

# 5. Quick rule-based decision
if range_df.max() / range_df.min() > 10:
    print("\n⚠ Feature scaling is likely necessary (large scale differences detected)")
else:
    print("\n✖ Feature scaling might not be strictly necessary (features on similar scale ranges)")

```

Summary Statistics:

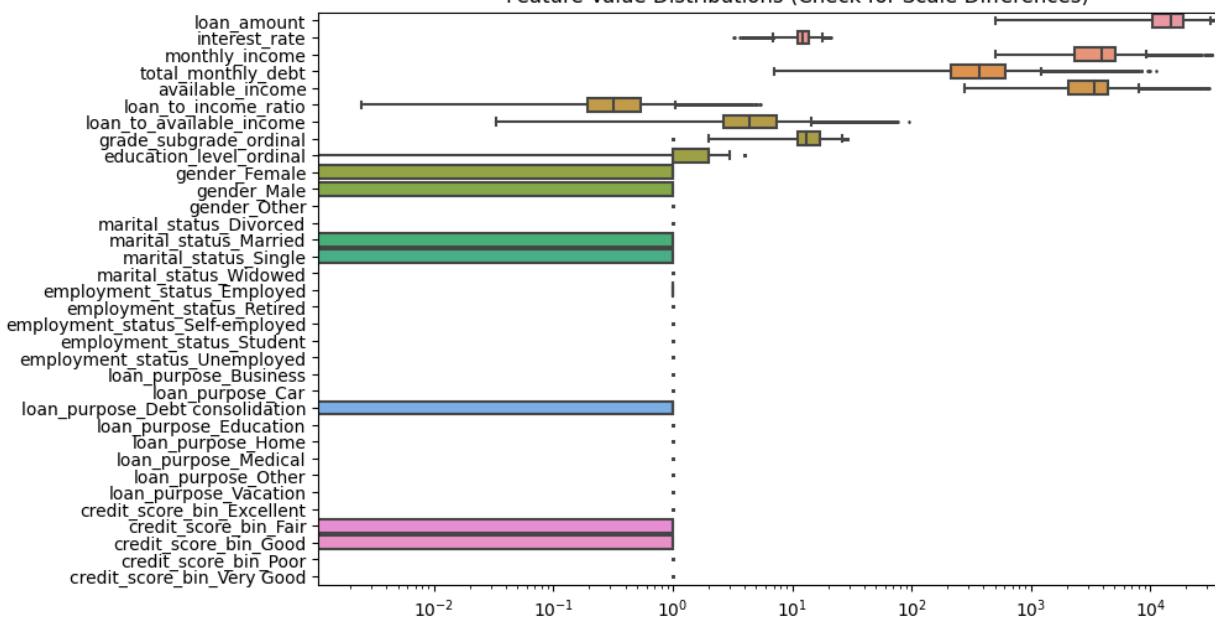
	loan_amount	interest_rate	monthly_income	total_monthly_debt	available_income
<b>count</b>	593994.000000	593994.000000	593994.000000	593994.000000	593994.000000
<b>mean</b>	15020.297629	12.356345	4017.683581	485.002319	3532.681262
<b>std</b>	6926.530568	2.008959	2225.995173	415.297292	1979.316733
<b>min</b>	500.090000	3.200000	500.202500	7.014620	274.692110
<b>25%</b>	10279.620000	10.990000	2327.866667	216.130683	2037.879680
<b>50%</b>	15000.220000	12.370000	3879.806667	367.691262	3386.199475
<b>75%</b>	18858.580000	13.680000	5081.776667	615.139751	4450.237361
<b>max</b>	48959.950000	20.990000	32781.811667	11151.413863	31477.877736

### Feature Ranges:

loan_amount	48459.860000
monthly_income	32281.609167
available_income	31203.185626
total_monthly_debt	11144.399243
loan_to_available_income	94.058593
grade_subgrade_ordinal	29.000000
interest_rate	17.790000
loan_to_income_ratio	5.458652
education_level_ordinal	4.000000
gender_Female	1.000000
gender_Male	1.000000
gender_Other	1.000000
marital_status_Divorced	1.000000
marital_status_Married	1.000000
marital_status_Single	1.000000
marital_status_Widowed	1.000000
employment_status_Employed	1.000000
employment_status_Retired	1.000000
employment_status_Self-employed	1.000000
employment_status_Student	1.000000
employment_status_Unemployed	1.000000
loan_purpose_Business	1.000000
loan_purpose_Car	1.000000
loan_purpose_Debt consolidation	1.000000
loan_purpose_Education	1.000000
loan_purpose_Home	1.000000
loan_purpose_Medical	1.000000
loan_purpose_Other	1.000000
loan_purpose_Vacation	1.000000
credit_score_bin_Excellent	1.000000
credit_score_bin_Fair	1.000000
credit_score_bin_Good	1.000000
credit_score_bin_Poor	1.000000
credit_score_bin_Very Good	1.000000

dtype: float64

Feature Value Distributions (Check for Scale Differences)



```
Features with significantly higher ranges: ['loan_amount', 'monthly_income',
'total_monthly_debt', 'available_income']
```

❖ Feature scaling is likely necessary (large scale differences detected).

## Train test split

```
In [17]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler, StandardScaler, RobustScaler, PowerTransformer
import numpy as np

# Use encoded data for model training
X = train_processed.drop("loan_paid_back", axis=1)
y = train_processed["loan_paid_back"]

# Ensure all columns are numeric
X = X.select_dtypes(include=[np.number])

# Train-test split
X_train, X_val, y_train, y_val = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# Choose scaling method
selected_method = 'Standard Scaling'

# Apply the best scaling method
if selected_method == 'Min-Max Scaling':
    scaler = MinMaxScaler()
elif selected_method == 'Standard Scaling':
    scaler = StandardScaler()
elif selected_method == 'Robust Scaling':
    scaler = RobustScaler()
elif selected_method == 'Power Transformation':
    scaler = PowerTransformer(method='yeo-johnson')
else:
    scaler = None

# Perform scaling
if scaler is not None:
    # Fit on training data
    X_train_scaled = scaler.fit_transform(X_train)
    # Only transform validation data
    X_val_scaled = scaler.transform(X_val)
else:
    X_train_scaled = X_train
    X_val_scaled = X_val

print("Scaling for tuning complete.")
```

Scaling for tuning complete.

```
In [18]: import optuna
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.metrics import log_loss, accuracy_score, roc_auc_score # <-- Corrected
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier

# ❖ Enable GPU for both XGBoost and LightGBM
```

```

# --- Hyperparameter tuning for XGBClassifier ---
def objective_xgb(trial):
    param = {
        'tree_method': 'gpu_hist',
        'predictor': 'gpu_predictor',
        'gpu_id': 0,
        'objective': 'binary:logistic', # <-- CORRECTED: Set objective for classification
        'lambda': trial.suggest_loguniform('lambda', 1e-3, 10.0),
        'alpha': trial.suggest_loguniform('alpha', 1e-3, 10.0),
        'colsample_bytree': trial.suggest_categorical('colsample_bytree', [0.3, 0.4, 0.5, 0.6, 0.7, 0.8]),
        'subsample': trial.suggest_categorical('subsample', [0.5, 0.6, 0.7, 0.8]),
        'learning_rate': trial.suggest_float('learning_rate', 0.005, 0.05, log=True),
        'n_estimators': trial.suggest_int('n_estimators', 200, 1000, step=100),
        'max_depth': trial.suggest_int('max_depth', 3, 12),
        'min_child_weight': trial.suggest_int('min_child_weight', 1, 300),
        'random_state': 42
    }

    model = XGBClassifier(**param, verbosity=0) # <-- CORRECTED: Classifier
    model.fit(X_train_scaled, y_train)
    y_pred_proba = model.predict_proba(X_val_scaled)[:, 1] # <-- Get probabilities
    ll = log_loss(y_val, y_pred_proba) # <-- CORRECTED: Use Log_Loss
    return ll

# --- Hyperparameter tuning for LGBMClassifier ---
def objective_lgbm(trial):
    param = {
        'device': 'gpu',
        'gpu_platform_id': 0,
        'gpu_device_id': 0,
        'boosting_type': 'gbdt',
        'objective': 'binary', # <-- CORRECTED: Set objective for classification
        'metric': 'binary_logloss', # <-- CORRECTED: Set metric
        'lambda_11': trial.suggest_float('lambda_11', 1e-5, 1.0, log=True),
        'lambda_12': trial.suggest_float('lambda_12', 1e-5, 1.0, log=True),
        'num_leaves': trial.suggest_int('num_leaves', 16, 256),
        'feature_fraction': trial.suggest_float('feature_fraction', 0.5, 1.0),
        'bagging_fraction': trial.suggest_float('bagging_fraction', 0.5, 1.0),
        'bagging_freq': trial.suggest_int('bagging_freq', 1, 7),
        'min_child_samples': trial.suggest_int('min_child_samples', 10, 100),
        'learning_rate': trial.suggest_float('learning_rate', 0.005, 0.3, log=True),
        'n_estimators': trial.suggest_int('n_estimators', 200, 1000, step=100),
        'max_depth': trial.suggest_int('max_depth', 3, 12),
        'random_state': 42,
        'verbosity': -1
    }

    model = LGBMClassifier(**param) # <-- CORRECTED: Classifier
    model.fit(X_train_scaled, y_train)
    y_pred_proba = model.predict_proba(X_val_scaled)[:, 1] # <-- Get probabilities
    ll = log_loss(y_val, y_pred_proba) # <-- CORRECTED: Use Log_Loss
    return ll

# --- Run GPU-accelerated Optuna optimization ---
print("⚡ Tuning XGBClassifier (GPU)...")
study_xgb = optuna.create_study(direction='minimize') # Minimize log_loss
study_xgb.optimize(objective_xgb, n_trials=50, timeout=600)
best_params_xgb = study_xgb.best_params
print(f"⚡ Best XGBClassifier parameters: {best_params_xgb}")

```

```

print("\n⚙️ Tuning LGBMClassifier (GPU)...")
study_lgbm = optuna.create_study(direction='minimize') # Minimize Log_Loss
study_lgbm.optimize(objective_lgbm, n_trials=50, timeout=600)
best_params_lgbm = study_lgbm.best_params
print(f"⚡️ Best LGBMClassifier parameters: {best_params_lgbm}")

# --- Initialize models with tuned GPU parameters ---
xgb_model = XGBClassifier(**best_params_xgb, tree_method='gpu_hist', predictor='gpu_hist')
lgbm_model = LGBMClassifier(**best_params_lgbm, device='gpu', objective='binary')

models = [
    ("XGBClassifier (GPU)", xgb_model),
    ("LGBMClassifier (GPU)", lgbm_model)
]

print("\n⚙️ Evaluating Tuned Models on GPU...\n")
logloss_scores = []
model_names = []

for name, model in models:
    model.fit(X_train_scaled, y_train)
    y_pred = model.predict(X_val_scaled)
    y_pred_proba = model.predict_proba(X_val_scaled)[:, 1]

    # --- CORRECTED Metrics ---
    ll = log_loss(y_val, y_pred_proba)
    acc = accuracy_score(y_val, y_pred)
    auc = roc_auc_score(y_val, y_pred_proba)

    logloss_scores.append(ll)
    model_names.append(name)

    print(f"{name:<30} | LogLoss: {ll:.5f} | Accuracy: {acc:.5f} | ROC AUC: {auc:.5f}")

```

[I 2025-11-03 07:15:12,165] A new study created in memory with name: no-name-8afcf609-3efc-4631-a3bc-8b843a3cd5f2

⚙️ Tuning XGBClassifier (GPU)...

```
[I 2025-11-03 07:15:17,864] Trial 0 finished with value: 0.2645231073170942 and parameters: {'lambda': 0.0361936500047033, 'alpha': 0.007568899354734025, 'colsample_bytree': 1.0, 'subsample': 0.6, 'learning_rate': 0.009507472541685866, 'n_estimators': 300, 'max_depth': 11, 'min_child_weight': 121}. Best is trial 0 with value: 0.2645231073170942.  
[I 2025-11-03 07:15:26,473] Trial 1 finished with value: 0.2573499981158369 and parameters: {'lambda': 0.09118575472550126, 'alpha': 0.006382495033513585, 'colsample_bytree': 0.8, 'subsample': 0.8, 'learning_rate': 0.020666723707219982, 'n_estimators': 900, 'max_depth': 8, 'min_child_weight': 58}. Best is trial 1 with value: 0.2573499981158369.  
[I 2025-11-03 07:15:32,359] Trial 2 finished with value: 0.2592545912919665 and parameters: {'lambda': 0.012995029534803772, 'alpha': 0.6651609154168596, 'colsample_bytree': 1.0, 'subsample': 0.7, 'learning_rate': 0.013830255474092062, 'n_estimators': 1000, 'max_depth': 4, 'min_child_weight': 202}. Best is trial 1 with value: 0.2573499981158369.  
[I 2025-11-03 07:15:37,303] Trial 3 finished with value: 0.2587124057571616 and parameters: {'lambda': 0.04113972357896702, 'alpha': 4.417408109316141, 'colsample_bytree': 0.7, 'subsample': 0.7, 'learning_rate': 0.025955877404414594, 'n_estimators': 400, 'max_depth': 11, 'min_child_weight': 218}. Best is trial 1 with value: 0.2573499981158369.  
[I 2025-11-03 07:15:40,227] Trial 4 finished with value: 0.2611889991191748 and parameters: {'lambda': 0.004734241446974353, 'alpha': 0.003013944723008013, 'colsample_bytree': 0.6, 'subsample': 0.6, 'learning_rate': 0.01710684720815652, 'n_estimators': 300, 'max_depth': 6, 'min_child_weight': 167}. Best is trial 1 with value: 0.2573499981158369.  
[I 2025-11-03 07:15:45,933] Trial 5 finished with value: 0.2576091697124722 and parameters: {'lambda': 0.12166336022200093, 'alpha': 0.014304734825826672, 'colsample_bytree': 0.7, 'subsample': 0.7, 'learning_rate': 0.0434660234370753, 'n_estimators': 600, 'max_depth': 1, 'min_child_weight': 91}. Best is trial 1 with value: 0.2573499981158369.  
[I 2025-11-03 07:15:53,635] Trial 6 finished with value: 0.2608557041858083 and parameters: {'lambda': 0.026980012268408095, 'alpha': 0.0018446481854280735, 'colsample_bytree': 0.7, 'subsample': 0.7, 'learning_rate': 0.013614857336067543, 'n_estimators': 700, 'max_depth': 9, 'min_child_weight': 185}. Best is trial 1 with value: 0.2573499981158369.  
[I 2025-11-03 07:15:58,081] Trial 7 finished with value: 0.26273495940873975 and parameters: {'lambda': 6.5082357772219535, 'alpha': 9.317028254420542, 'colsample_bytree': 1.0, 'subsample': 0.6, 'learning_rate': 0.009085498824438855, 'n_estimators': 700, 'max_depth': 4, 'min_child_weight': 64}. Best is trial 1 with value: 0.2573499981158369.  
[I 2025-11-03 07:16:05,311] Trial 8 finished with value: 0.2579580561236633 and parameters: {'lambda': 0.003380654226157312, 'alpha': 0.01703146881244035, 'colsample_bytree': 0.7, 'subsample': 1.0, 'learning_rate': 0.04202674720091878, 'n_estimators': 700, 'max_depth': 10, 'min_child_weight': 286}. Best is trial 1 with value: 0.2573499981158369.  
[I 2025-11-03 07:16:10,779] Trial 9 finished with value: 0.25816211054741106 and parameters: {'lambda': 0.0031406438367416823, 'alpha': 0.12918290371363503, 'colsample_bytree': 0.9, 'subsample': 0.7, 'learning_rate': 0.024981132832192814, 'n_estimators': 800, 'max_depth': 5, 'min_child_weight': 181}. Best is trial 1 with value: 0.2573499981158369.  
[I 2025-11-03 07:16:20,395] Trial 10 finished with value: 0.2594719159297115 and parameters: {'lambda': 1.5661425752614788, 'alpha': 0.07643494139376109, 'colsample_bytree': 0.7, 'subsample': 0.8, 'learning_rate': 0.005489057691074062, 'n_estimators': 1000, 'max_depth': 7, 'min_child_weight': 5}. Best is trial 1 with value: 0.2573499981158369.  
[I 2025-11-03 07:16:25,532] Trial 11 finished with value: 0.2577697358008288 and parameters: {'lambda': 0.24326017984851894, 'alpha': 0.025526427065941656, 'colsample_bytree': 0.5, 'subsample': 0.5, 'learning_rate': 0.04906657470555916, 'n_estimators': 500, 'max_depth': 8, 'min_child_weight': 90}. Best is trial 1 with value: 0.2573499981158369.  
[I 2025-11-03 07:16:34,323] Trial 12 finished with value: 0.25750789239913274 and parameters: {'lambda': 0.24133162318144, 'alpha': 0.0010088575703394705, 'colsample_bytree': 0.7, 'subsample': 0.8, 'learning_rate': 0.030039663735225014, 'n_estimators': 900, 'max_depth': 8, 'min_child_weight': 32}. Best is trial 1 with value: 0.2573499981158369.  
[I 2025-11-03 07:16:42,472] Trial 13 finished with value: 0.2574224778365216 and parameters: {'lambda': 0.5266193908344308, 'alpha': 0.0014151126501476, 'colsample_bytree': 0.7, 'subsample': 0.8, 'learning_rate': 0.027853433914135396, 'n_estimators': 900, 'max_depth': 7, 'min_child_weight': 2}. Best is trial 1 with value: 0.2573499981158369.
```

```
[I 2025-11-03 07:16:49,652] Trial 14 finished with value: 0.25754753481300513 and
parameters: {'lambda': 1.0566261705611366, 'alpha': 0.004510060596708759,
'colsample_bytree': 0.7, 'subsample': 0.8, 'learning_rate': 0.018133034864827927,
'n_estimators': 900, 'max_depth': 6, 'min_child_weight': 13}. Best is trial 1 with va
0.2573499981158369.
[I 2025-11-03 07:17:02,748] Trial 15 finished with value: 0.2587000782480725 and
parameters: {'lambda': 0.9219592386752085, 'alpha': 0.07262882153129689,
'colsample_bytree': 0.7, 'subsample': 0.8, 'learning_rate': 0.03258283887759407,
'n_estimators': 900, 'max_depth': 12, 'min_child_weight': 50}. Best is trial 1 with v
0.2573499981158369.
[I 2025-11-03 07:17:09,232] Trial 16 finished with value: 0.2578573695826229 and
parameters: {'lambda': 5.095116237085027, 'alpha': 0.33996098929077445, 'colsample_b
0.5, 'subsample': 0.8, 'learning_rate': 0.019407058782810696, 'n_estimators': 800,
'max_depth': 6, 'min_child_weight': 123}. Best is trial 1 with value: 0.2573499981158
[I 2025-11-03 07:17:14,771] Trial 17 finished with value: 0.2585088610381899 and
parameters: {'lambda': 0.42624188613345126, 'alpha': 0.0010980232948895675,
'colsample_bytree': 0.7, 'subsample': 0.5, 'learning_rate': 0.02216773352409775,
'n_estimators': 1000, 'max_depth': 3, 'min_child_weight': 1}. Best is trial 1 with va
0.2573499981158369.
[I 2025-11-03 07:17:24,777] Trial 18 finished with value: 0.25798091772046844 and
parameters: {'lambda': 0.0809909688327199, 'alpha': 0.005547614572827067,
'colsample_bytree': 0.7, 'subsample': 1.0, 'learning_rate': 0.010389433504091137,
'n_estimators': 800, 'max_depth': 9, 'min_child_weight': 73}. Best is trial 1 with va
0.2573499981158369.
[I 2025-11-03 07:17:30,364] Trial 19 finished with value: 0.25731562046172163 and
parameters: {'lambda': 2.746270697088073, 'alpha': 0.03333123632829805, 'colsample_b
0.7, 'subsample': 0.8, 'learning_rate': 0.03643335365651437, 'n_estimators': 600,
'max_depth': 7, 'min_child_weight': 37}. Best is trial 19 with value: 0.2573156204617
[I 2025-11-03 07:17:36,278] Trial 20 finished with value: 0.2576358598043774 and
parameters: {'lambda': 0.010053749681438472, 'alpha': 0.03977407477009539,
'colsample_bytree': 0.3, 'subsample': 0.8, 'learning_rate': 0.03574734490967888,
'n_estimators': 500, 'max_depth': 9, 'min_child_weight': 126}. Best is trial 19 with
0.25731562046172163.
[I 2025-11-03 07:17:41,886] Trial 21 finished with value: 0.2574151157221354 and
parameters: {'lambda': 2.8584161451684267, 'alpha': 0.00891615835995151,
'colsample_bytree': 0.7, 'subsample': 0.8, 'learning_rate': 0.02896741425641675,
'n_estimators': 600, 'max_depth': 7, 'min_child_weight': 44}. Best is trial 19 with
0.25731562046172163.
[I 2025-11-03 07:17:47,481] Trial 22 finished with value: 0.2573876680389996 and
parameters: {'lambda': 3.345061080070869, 'alpha': 0.010009968495511242,
'colsample_bytree': 0.7, 'subsample': 0.8, 'learning_rate': 0.03449612395708449,
'n_estimators': 600, 'max_depth': 7, 'min_child_weight': 37}. Best is trial 19 with
0.25731562046172163.
[I 2025-11-03 07:17:51,355] Trial 23 finished with value: 0.25781124024494334 and
parameters: {'lambda': 9.669034384991654, 'alpha': 0.2190456458994139, 'colsample_b
0.7, 'subsample': 0.8, 'learning_rate': 0.038297716415108665, 'n_estimators': 500,
'max_depth': 5, 'min_child_weight': 35}. Best is trial 19 with value: 0.2573156204617
[I 2025-11-03 07:17:56,932] Trial 24 finished with value: 0.25778624646035775 and
parameters: {'lambda': 2.6037127357009724, 'alpha': 0.0337677110336313, 'colsample_b
0.7, 'subsample': 0.8, 'learning_rate': 0.020751912960964312, 'n_estimators': 600,
'max_depth': 7, 'min_child_weight': 87}. Best is trial 19 with value: 0.2573156204617
[I 2025-11-03 07:18:00,451] Trial 25 finished with value: 0.2589067255803117 and
parameters: {'lambda': 0.001044885706690124, 'alpha': 0.010674149868349192,
'colsample_bytree': 0.5, 'subsample': 0.5, 'learning_rate': 0.035043824089197415,
'n_estimators': 200, 'max_depth': 10, 'min_child_weight': 59}. Best is trial 19 with
0.25731562046172163.
[I 2025-11-03 07:18:04,105] Trial 26 finished with value: 0.25757803826289094 and
parameters: {'lambda': 0.08264939587100727, 'alpha': 0.003181821772459335,
'colsample_bytree': 0.7, 'subsample': 0.8, 'learning_rate': 0.04948259538542322,
'n_estimators': 400, 'max_depth': 6, 'min_child_weight': 26}. Best is trial 19 with
```

0.25731562046172163.  
[I 2025-11-03 07:18:09,906] Trial 27 finished with value: 0.25828610667991514 and parameters: {'lambda': 3.073690481478437, 'alpha': 1.259177716837584, 'colsample\_bytree': 0.7, 'subsample': 1.0, 'learning\_rate': 0.023023357426613246, 'n\_estimators': 600, 'max\_depth': 8, 'min\_child\_weight': 246}. Best is trial 19 with value: 0.25731562046172163.  
[I 2025-11-03 07:18:15,085] Trial 28 finished with value: 0.2587965005045791 and parameters: {'lambda': 0.19783695725050118, 'alpha': 0.03998320452885442, 'colsample\_bytree': 0.7, 'subsample': 0.8, 'learning\_rate': 0.01486421346139551, 'n\_estimators': 700, 'max\_depth': 5, 'min\_child\_weight': 111}. Best is trial 19 with value: 0.25731562046172163.  
[I 2025-11-03 07:18:20,400] Trial 29 finished with value: 0.2598700920807726 and parameters: {'lambda': 0.5333715889110444, 'alpha': 0.008349225767380257, 'colsample\_bytree': 1.0, 'subsample': 0.6, 'learning\_rate': 0.011817633530455503, 'n\_estimators': 400, 'max\_depth': 10, 'min\_child\_weight': 141}. Best is trial 19 with value: 0.25731562046172163.  
[I 2025-11-03 07:18:27,596] Trial 30 finished with value: 0.27193903588524276 and parameters: {'lambda': 0.03497097150162013, 'alpha': 0.020900653667307725, 'colsample\_bytree': 0.5, 'subsample': 0.8, 'learning\_rate': 0.006343452402425456, 'n\_estimators': 500, 'max\_depth': 9, 'min\_child\_weight': 100}. Best is trial 19 with value: 0.25731562046172163.  
[I 2025-11-03 07:18:33,213] Trial 31 finished with value: 0.25740248782501746 and parameters: {'lambda': 3.562625192591287, 'alpha': 0.007805314110200777, 'colsample\_bytree': 0.7, 'subsample': 0.8, 'learning\_rate': 0.028884255018318167, 'n\_estimators': 600, 'max\_depth': 7, 'min\_child\_weight': 44}. Best is trial 19 with value: 0.25731562046172163.  
[I 2025-11-03 07:18:39,626] Trial 32 finished with value: 0.25739958318073713 and parameters: {'lambda': 1.39482674157152, 'alpha': 0.0028708926724178496, 'colsample\_bytree': 0.7, 'subsample': 0.8, 'learning\_rate': 0.03190897995172127, 'n\_estimators': 700, 'max\_depth': 7, 'min\_child\_weight': 24}. Best is trial 19 with value: 0.25731562046172163.  
[I 2025-11-03 07:18:47,071] Trial 33 finished with value: 0.2575988145313977 and parameters: {'lambda': 1.493698394869845, 'alpha': 0.0026893895290535097, 'colsample\_bytree': 1.0, 'subsample': 0.8, 'learning\_rate': 0.04217122951641485, 'n\_estimators': 800, 'max\_depth': 8, 'min\_child\_weight': 72}. Best is trial 19 with value: 0.25731562046172163.  
[I 2025-11-03 07:18:52,695] Trial 34 finished with value: 0.25737800681748224 and parameters: {'lambda': 1.7132032645818234, 'alpha': 0.0053424955562438, 'colsample\_bytree': 0.7, 'subsample': 0.8, 'learning\_rate': 0.03252593950740572, 'n\_estimators': 700, 'max\_depth': 6, 'min\_child\_weight': 30}. Best is trial 19 with value: 0.25731562046172163.  
[I 2025-11-03 07:18:57,073] Trial 35 finished with value: 0.25808442532015785 and parameters: {'lambda': 9.391317787021466, 'alpha': 0.005218083909051403, 'colsample\_bytree': 0.9, 'subsample': 0.6, 'learning\_rate': 0.024241465833981526, 'n\_estimators': 500, 'max\_depth': 6, 'min\_child\_weight': 23}. Best is trial 19 with value: 0.25731562046172163.  
[I 2025-11-03 07:18:59,557] Trial 36 finished with value: 0.25916812554118107 and parameters: {'lambda': 4.909784171471608, 'alpha': 0.013490854289304358, 'colsample\_bytree': 0.7, 'subsample': 0.8, 'learning\_rate': 0.03953841169880165, 'n\_estimators': 300, 'max\_depth': 4, 'min\_child\_weight': 75}. Best is trial 19 with value: 0.25731562046172163.  
[I 2025-11-03 07:19:05,344] Trial 37 finished with value: 0.25924028289561596 and parameters: {'lambda': 0.05715600687294148, 'alpha': 0.05806199432114279, 'colsample\_bytree': 0.3, 'subsample': 0.7, 'learning\_rate': 0.01730464958553354, 'n\_estimators': 800, 'max\_depth': 5, 'min\_child\_weight': 54}. Best is trial 19 with value: 0.25731562046172163.  
[I 2025-11-03 07:19:10,693] Trial 38 finished with value: 0.2582211941306502 and parameters: {'lambda': 0.01549896014674224, 'alpha': 0.018020943075521705, 'colsample\_bytree': 0.9, 'subsample': 1.0, 'learning\_rate': 0.034835410475615995, 'n\_estimators': 700, 'max\_depth': 6, 'min\_child\_weight': 223}. Best is trial 19 with value: 0.25731562046172163.  
[I 2025-11-03 07:19:16,408] Trial 39 finished with value: 0.25824831990058866 and

```
parameters: {'lambda': 0.879383651614276, 'alpha': 0.005041122652933428, 'colsample_bytree': 1.0, 'subsample': 0.5, 'learning_rate': 0.026443544364382947, 'n_estimators': 600, 'max_depth': 8, 'min_child_weight': 145}. Best is trial 19 with 0.25731562046172163.  
[I 2025-11-03 07:19:22,063] Trial 40 finished with value: 0.25758404343100055 and parameters: {'lambda': 0.12084691734055211, 'alpha': 0.0021322887924655498, 'colsample_bytree': 0.7, 'subsample': 0.6, 'learning_rate': 0.043627974503275185, 'n_estimators': 700, 'max_depth': 6, 'min_child_weight': 14}. Best is trial 19 with 0.25731562046172163.  
[I 2025-11-03 07:19:28,417] Trial 41 finished with value: 0.25740751772889653 and parameters: {'lambda': 1.53067226017384, 'alpha': 0.003228818747108632, 'colsample_bytree': 0.7, 'subsample': 0.8, 'learning_rate': 0.031779294830742454, 'n_estimators': 700, 'max_depth': 7, 'min_child_weight': 36}. Best is trial 19 with value: 0.25731562046172163.  
[I 2025-11-03 07:19:34,907] Trial 42 finished with value: 0.2573833416910459 and parameters: {'lambda': 1.8903837220166977, 'alpha': 0.001821357211682189, 'colsample_bytree': 0.7, 'subsample': 0.8, 'learning_rate': 0.032855631974691185, 'n_estimators': 700, 'max_depth': 7, 'min_child_weight': 20}. Best is trial 19 with 0.25731562046172163.  
[I 2025-11-03 07:19:40,871] Trial 43 finished with value: 0.2575609414293694 and parameters: {'lambda': 2.197432379630974, 'alpha': 0.0017090459750847516, 'colsample_bytree': 0.7, 'subsample': 0.8, 'learning_rate': 0.04580523617648304, 'n_estimators': 600, 'max_depth': 8, 'min_child_weight': 61}. Best is trial 19 with 0.25731562046172163.  
[I 2025-11-03 07:19:48,157] Trial 44 finished with value: 0.2576033590333954 and parameters: {'lambda': 0.39895164471356576, 'alpha': 0.012248725855203401, 'colsample_bytree': 0.7, 'subsample': 0.7, 'learning_rate': 0.03898593799710853, 'n_estimators': 800, 'max_depth': 7, 'min_child_weight': 15}. Best is trial 19 with 0.25731562046172163.  
[I 2025-11-03 07:19:57,175] Trial 45 finished with value: 0.2570974993002868 and parameters: {'lambda': 5.886232957483388, 'alpha': 0.025940388011077873, 'colsample_bytree': 0.3, 'subsample': 0.8, 'learning_rate': 0.026569438393859245, 'n_estimators': 900, 'max_depth': 8, 'min_child_weight': 46}. Best is trial 45 with 0.2570974993002868.  
[I 2025-11-03 07:20:06,653] Trial 46 finished with value: 0.2572943246155355 and parameters: {'lambda': 7.553103188892479, 'alpha': 0.13438983425704323, 'colsample_bytree': 0.3, 'subsample': 0.8, 'learning_rate': 0.02615716099718383, 'n_estimators': 900, 'max_depth': 9, 'min_child_weight': 100}. Best is trial 45 with value: 0.2570974993002868.  
[I 2025-11-03 07:20:18,265] Trial 47 finished with value: 0.25734227315159125 and parameters: {'lambda': 5.8455752130841905, 'alpha': 0.1231452832023838, 'colsample_bytree': 0.3, 'subsample': 0.8, 'learning_rate': 0.020840511602705575, 'n_estimators': 1000, 'max_depth': 10, 'min_child_weight': 101}. Best is trial 45 with value: 0.2570974993002868.  
[I 2025-11-03 07:20:30,917] Trial 48 finished with value: 0.2574147253434204 and parameters: {'lambda': 6.3465793164448319, 'alpha': 0.13384077799292762, 'colsample_bytree': 0.3, 'subsample': 0.8, 'learning_rate': 0.020161665431804185, 'n_estimators': 1000, 'max_depth': 11, 'min_child_weight': 101}. Best is trial 45 with value: 0.2570974993002868.  
[I 2025-11-03 07:20:41,257] Trial 49 finished with value: 0.25833729942246175 and parameters: {'lambda': 5.240139445349718, 'alpha': 0.7204094288481924, 'colsample_bytree': 0.3, 'subsample': 0.7, 'learning_rate': 0.015534267841260203, 'n_estimators': 1000, 'max_depth': 9, 'min_child_weight': 170}. Best is trial 45 with value: 0.2570974993002868.  
[I 2025-11-03 07:20:41,258] A new study created in memory with name: no-name-24d9e47d-39ff-4841-9e2b-2c86755c9e2a
```

✓ Best XGBClassifier parameters: {'lambda': 5.886232957483388, 'alpha': 0.025940388011077873, 'colsample\_bytree': 0.3, 'subsample': 0.8, 'learning\_rate': 0.026569438393859245, 'n\_estimators': 900, 'max\_depth': 8, 'min\_child\_weight': 46}

⌚ Tuning LGBMClassifier (GPU)...



```
[I 2025-11-03 07:23:28,097] Trial 5 finished with value: 0.25771091148873165 and parameters: {'lambda_11': 0.015182241770413525, 'lambda_12': 0.0015200752229676073, 'num_leaves': 246, 'feature_fraction': 0.8880621678379127, 'bagging_fraction': 0.6964318809356786, 'bagging_freq': 5, 'min_child_samples': 80, 'learning_rate': 0.1749633527686237, 'n_estimators': 800, 'max_depth': 3}. Best is trial 3 with value: 0.25743496045312714.  
[I 2025-11-03 07:23:39,098] Trial 6 finished with value: 0.2581490067176458 and parameters: {'lambda_11': 0.26673972807827945, 'lambda_12': 0.0054649190527836485, 'num_leaves': 246, 'feature_fraction': 0.6762239561642109, 'bagging_fraction': 0.8174726291630314, 'bagging_freq': 7, 'min_child_samples': 32, 'learning_rate': 0.0770329904540322, 'n_estimators': 400, 'max_depth': 3}. Best is trial 3 with value: 0.25743496045312714.  
[I 2025-11-03 07:24:04,648] Trial 7 finished with value: 0.2616319007377422 and parameters: {'lambda_11': 0.0245656669975974, 'lambda_12': 0.0032108559885385577, 'num_leaves': 246, 'feature_fraction': 0.8898431880657955, 'bagging_fraction': 0.7350873042996415, 'bagging_freq': 5, 'min_child_samples': 93, 'learning_rate': 0.08686248711625927, 'n_estimators': 400, 'max_depth': 12}. Best is trial 3 with value: 0.25743496045312714.  
[I 2025-11-03 07:24:27,728] Trial 8 finished with value: 0.2585985096333796 and parameters: {'lambda_11': 0.0027781886246329474, 'lambda_12': 0.00037221988853414494, 'num_leaves': 249, 'feature_fraction': 0.8363847875841872, 'bagging_fraction': 0.6422222655715286, 'bagging_freq': 1, 'min_child_samples': 93, 'learning_rate': 0.012170551574233664, 'n_estimators': 800, 'max_depth': 5}. Best is trial 3 with value: 0.25743496045312714.  
[I 2025-11-03 07:24:41,252] Trial 9 finished with value: 0.26294327825373964 and parameters: {'lambda_11': 0.857595416585917, 'lambda_12': 0.00018306083147395157, 'num_leaves': 178, 'feature_fraction': 0.7276619708739973, 'bagging_fraction': 0.6314031142147365, 'bagging_freq': 2, 'min_child_samples': 81, 'learning_rate': 0.01549029823847191, 'n_estimators': 300, 'max_depth': 5}. Best is trial 3 with value: 0.25743496045312714.  
[I 2025-11-03 07:25:27,399] Trial 10 finished with value: 0.26119359790080376 and parameters: {'lambda_11': 2.0925040665937483e-05, 'lambda_12': 1.5476848641847573e-05, 'num_leaves': 21, 'feature_fraction': 0.5396067837676868, 'bagging_fraction': 0.8506396331888332, 'bagging_freq': 6, 'min_child_samples': 52, 'learning_rate': 0.006048999174428362, 'n_estimators': 1000, 'max_depth': 9}. Best is trial 3 with value: 0.25743496045312714.  
[I 2025-11-03 07:26:21,918] Trial 11 finished with value: 0.2582849761729365 and parameters: {'lambda_11': 0.00014127218084153097, 'lambda_12': 0.054740335783379945, 'num_leaves': 126, 'feature_fraction': 0.7879224442292634, 'bagging_fraction': 0.9691918108332408, 'bagging_freq': 7, 'min_child_samples': 61, 'learning_rate': 0.0345184775996353, 'n_estimators': 1000, 'max_depth': 12}. Best is trial 3 with value: 0.25743496045312714.  
[I 2025-11-03 07:27:22,171] Trial 12 finished with value: 0.2596104828034562 and parameters: {'lambda_11': 0.00026919936857296015, 'lambda_12': 0.015559082245543743, 'num_leaves': 98, 'feature_fraction': 0.5681303773228598, 'bagging_fraction': 0.9799202441381428, 'bagging_freq': 6, 'min_child_samples': 69, 'learning_rate': 0.006497042299170708, 'n_estimators': 800, 'max_depth': 10}. Best is trial 3 with value: 0.25743496045312714.  
[I 2025-11-03 07:28:11,850] Trial 13 finished with value: 0.25725828269692314 and parameters: {'lambda_11': 0.0002940013123952954, 'lambda_12': 4.5225371776859564e-05, 'num_leaves': 207, 'feature_fraction': 0.6462374794303218, 'bagging_fraction': 0.8567269768323934, 'bagging_freq': 7, 'min_child_samples': 42, 'learning_rate': 0.020079871353289785, 'n_estimators': 900, 'max_depth': 7}. Best is trial 13 with value: 0.25725828269692314.  
[I 2025-11-03 07:28:44,157] Trial 14 finished with value: 0.2574534143995856 and parameters: {'lambda_11': 1.3517869679106238e-05, 'lambda_12': 1.5632886280360352e-05, 'num_leaves': 229, 'feature_fraction': 0.606454224560812, 'bagging_fraction': 0.857848207904601, 'bagging_freq': 6, 'min_child_samples': 41, 'learning_rate': 0.039045441981621194, 'n_estimators': 600, 'max_depth': 7}. Best is trial 13 with value: 0.25725828269692314.  
[I 2025-11-03 07:29:28,154] Trial 15 finished with value: 0.2642831285010989 and parameters: {'lambda_11': 7.3516609299707e-05, 'lambda_12': 5.248591374988818e-05, 'num_leaves': 201, 'feature_fraction': 0.7692467630158706, 'bagging_fraction':
```

```

0.7879759882575366, 'bagging_freq': 5, 'min_child_samples': 30, 'learning_rate': 0.10412394927042785, 'n_estimators': 700, 'max_depth': 8}. Best is trial 13 with value 0.25725828269692314.
[I 2025-11-03 07:30:11,957] Trial 16 finished with value: 0.25723239986063656 and parameters: {'lambda_l1': 0.0006130434418029652, 'lambda_l2': 6.376572790581942e-05, 'num_leaves': 107, 'feature_fraction': 0.6120751671862034, 'bagging_fraction': 0.8940030237623764, 'bagging_freq': 4, 'min_child_samples': 45, 'learning_rate': 0.03996637371759323, 'n_estimators': 900, 'max_depth': 6}. Best is trial 16 with value 0.25723239986063656.
[I 2025-11-03 07:31:02,547] Trial 17 finished with value: 0.25710944888518794 and parameters: {'lambda_l1': 0.0006676400275768641, 'lambda_l2': 6.882280051749269e-05, 'num_leaves': 155, 'feature_fraction': 0.5063358344178233, 'bagging_fraction': 0.9060032142363916, 'bagging_freq': 3, 'min_child_samples': 43, 'learning_rate': 0.026077518808694123, 'n_estimators': 700, 'max_depth': 8}. Best is trial 17 with value 0.25710944888518794.
▽ Best LGBMClassifier parameters: {'lambda_l1': 0.0006676400275768641, 'lambda_l2': 6.882280051749269e-05, 'num_leaves': 155, 'feature_fraction': 0.5063358344178233, 'bagging_fraction': 0.9060032142363916, 'bagging_freq': 3, 'min_child_samples': 43, 'learning_rate': 0.026077518808694123, 'n_estimators': 700, 'max_depth': 8}

```

⚡ Evaluating Tuned Models on GPU...

XGBClassifier (GPU)	LogLoss: 0.25703   Accuracy: 0.90212   ROC AUC: 0.9:
LGBMClassifier (GPU)	LogLoss: 0.25734   Accuracy: 0.90231   ROC AUC: 0.9:

```
In [19]: # Select best model
best_idx = np.argmin(logloss_scores) # <-- Select based on Lowest Log_Loss
best_model_name = model_names[best_idx]
best_model = models[best_idx][1]
print(f"\n▽ Best Model Based on LogLoss: {best_model_name}")

# Evaluate final model
y_pred = best_model.predict(X_val_scaled)
y_pred_proba = best_model.predict_proba(X_val_scaled)[:, 1]

ll_final = log_loss(y_val, y_pred_proba)
acc_final = accuracy_score(y_val, y_pred)
auc_final = roc_auc_score(y_val, y_pred_proba)

print("\n▽ Final Model Evaluation (on validation set):")
print(f"LogLoss : {ll_final:.5f}")
print(f"Accuracy : {acc_final:.5f}")
print(f"ROC AUC Score: {auc_final:.5f}")
```

▽ Best Model Based on LogLoss: XGBClassifier (GPU)

▽ Final Model Evaluation (on validation set):  
 LogLoss : 0.25703  
 Accuracy : 0.90212  
 ROC AUC Score: 0.91283

## Selecting best model and Generating Submission

```
In [20]: print("\n⚡ Retraining the best model on full training data...")

# Prepare full training features and target
X_full = train_processed.drop(columns=['loan_paid_back'], errors='ignore')
y_full = train_processed['loan_paid_back']

# Ensure all columns are numeric
```

```

X_full = X_full.select_dtypes(include=[np.number])

# Scale full data using the same scaler
if scaler is not None:
    # We must re-fit the scaler on the FULL training data
    X_full_scaled = scaler.fit_transform(X_full)
else:
    X_full_scaled = X_full

# Retrain best model on the full scaled dataset
best_model.fit(X_full_scaled, y_full)

print(f"☛ Model retrained successfully: {best_model_name} ({best_model.__class__}")

```

☛ Retraining the best model on full training data...  
 ☻ Model retrained successfully: XGBClassifier (GPU) (XGBClassifier)

## Submission creation

In [21]:

```

X_submission = test_processed.select_dtypes(include=[np.number])

# Ensure column order matches the training data
X_submission = X_submission[X_full.columns]

# Scale using the scaler that was fit on X_full
if scaler is not None:
    X_submission_scaled = scaler.transform(X_submission)
else:
    X_submission_scaled = X_submission

```

In [22]:

```

print("⌚ Generating predictions using the best model...")
# Predict probabilities for the submission
submission_preds = best_model.predict_proba(X_submission_scaled)[:, 1]

# Optional: clip predictions to valid range [0, 1]
submission_preds = np.clip(submission_preds, 0, 1)

```

⌚ Generating predictions using the best model...

In [23]:

```

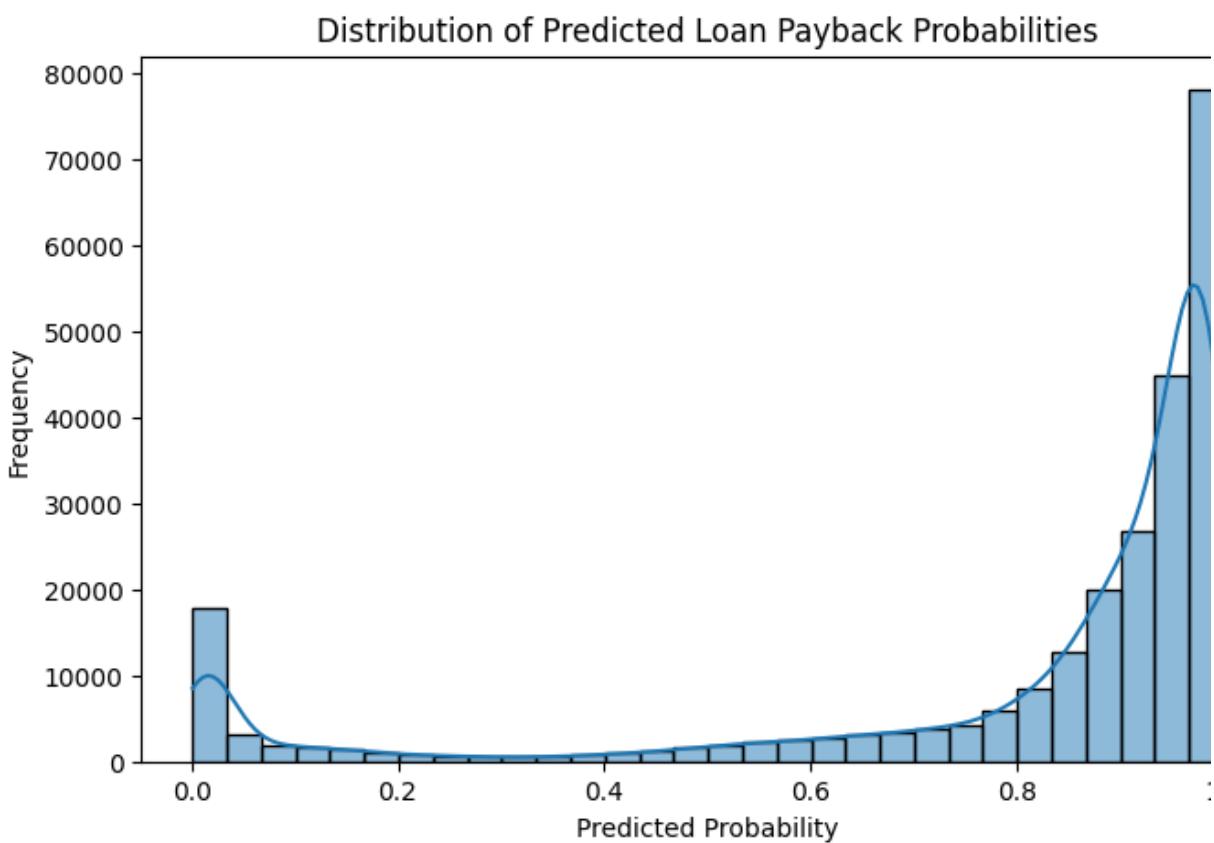
submission = pd.DataFrame({
    'id': test_ids,
    'loan_paid_back': submission_preds
})
submission.to_csv('submission.csv', index=False)
print("\n☛ Submission file 'submission.csv' generated successfully!")
display(submission.head())

```

☛ Submission file 'submission.csv' generated successfully!

	<b>id</b>	<b>loan_paid_back</b>
<b>0</b>	593994	0.944431
<b>1</b>	593995	0.955119
<b>2</b>	593996	0.431062
<b>3</b>	593997	0.962783
<b>4</b>	593998	0.978027

```
In [24]: plt.figure(figsize=(8, 5))
sns.histplot(submission['loan_paid_back'], bins=30, kde=True)
plt.title('Distribution of Predicted Loan Payback Probabilities')
plt.xlabel('Predicted Probability')
plt.ylabel('Frequency')
plt.show()
```



```
In [ ]:
```