

Kaggle Playground

Problem Statement / Real World Implementations

```
In [1]: # --- 1. Importing Libraries ---
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
from scipy.stats import rankdata

# Preprocessing
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import roc_auc_score

# Models
import lightgbm as lgb
import xgboost as xgb
from catboost import CatBoostClassifier

# Notebook settings
warnings.filterwarnings('ignore')
pd.set_option('display.max_columns', None)
```

```
In [2]: # --- 2. Configuration ---
class Config:
    """Configuration class for hyperparameters and settings"""
    N_SPLITS = 5
    SEED = 42
    TARGET = 'loan_paid_back'

config = Config()

print("Version 11: 3-Model CV Ensemble (LGBM+XGB+CAT) with Advanced FE")
print(f"--- 1. Loading Data ---")
```

```
Version 11: 3-Model CV Ensemble (LGBM+XGB+CAT) with Advanced FE
--- 1. Loading Data ---
```

```
In [3]: # Define file paths
TRAIN_PATH = "/kaggle/input/playground-series-s5e11/train.csv"
TEST_PATH = "/kaggle/input/playground-series-s5e11/test.csv"
SUBMISSION_PATH = "/kaggle/input/playground-series-s5e11/sample_submission.csv"

# Load the datasets
train = pd.read_csv(TRAIN_PATH)
test = pd.read_csv(TEST_PATH)
sample_submission = pd.read_csv(SUBMISSION_PATH)
```

```
print(f"Train shape: {train.shape}, Test shape: {test.shape}")
print("--- 2. Defining Preprocessing & Feature Engineering ---")
```

Train shape: (593994, 13), Test shape: (254569, 12)
--- 2. Defining Preprocessing & Feature Engineering ---

```
In [4]: def complete_feature_engineering(df):
    """
    Comprehensive feature engineering pipeline for loan prediction
    """
    df = df.copy()

    # 1. FINANCIAL RATIOS
    df['loan_to_income_ratio'] = df['loan_amount'] / (df['annual_income'] + 1)
    df['monthly_income'] = df['annual_income'] / 12
    # Simplified approximation from source
    df['monthly_payment_estimate'] = (df['loan_amount'] * df['interest_rate']) / 10
    df['payment_to_income_ratio'] = df['monthly_payment_estimate'] / (df['monthly_income'])
    df['current_debt_amount'] = df['debt_to_income_ratio'] * df['annual_income']
    df['total_debt_with_loan'] = df['current_debt_amount'] + df['loan_amount']
    df['new_debt_to_income'] = df['total_debt_with_loan'] / (df['annual_income'] + 1)
    df['debt_increase_ratio'] = df['new_debt_to_income'] / (df['debt_to_income_ratio'])
    df['disposable_income'] = df['annual_income'] - df['current_debt_amount']
    df['disposable_income_ratio'] = df['disposable_income'] / (df['annual_income'])
    df['loan_to_disposable_income'] = df['loan_amount'] / (df['disposable_income'])
    df['monthly_disposable_income'] = df['disposable_income'] / 12
    df['payment_to_disposable_ratio'] = df['monthly_payment_estimate'] / (df['monthly_disposable_income'])
    df['annual_payment_burden'] = df['monthly_payment_estimate'] * 12
    df['payment_burden_ratio'] = df['annual_payment_burden'] / (df['annual_income'])

    # 2. CREDIT SCORE FEATURES
    df['credit_score_normalized'] = df['credit_score'] / 850
    df['credit_risk_score'] = 1 - df['credit_score_normalized']
    df['credit_score_squared'] = df['credit_score'] ** 2
    df['credit_score_log'] = np.log1p(df['credit_score'])
    df['credit_category'] = pd.cut(df['credit_score'], bins=[0, 580, 670, 740, 800],
                                    labels=['poor', 'fair', 'good', 'very_good', 'excellent'])
    df['credit_income_interaction'] = df['credit_score'] * df['annual_income']
    df['credit_times_dti'] = df['credit_score'] * df['debt_to_income_ratio']
    df['credit_loan_interaction'] = df['credit_score'] * df['loan_amount']

    # 3. INTEREST RATE FEATURES
    df['high_interest_flag'] = (df['interest_rate'] > df['interest_rate'].median())
    df['very_high_interest'] = (df['interest_rate'] > df['interest_rate'].quantile(0.9))
    df['low_interest_flag'] = (df['interest_rate'] < df['interest_rate'].quantile(0.1))
    df['total_interest_cost'] = df['loan_amount'] * df['interest_rate'] / 100
    df['interest_burden'] = df['total_interest_cost'] / (df['annual_income'] + 1)
    df['interest_credit_mismatch'] = df['interest_rate'] * (1 - df['credit_score_normalized'])
    df['interest_credit_ratio'] = df['interest_rate'] / (df['credit_score'] / 100)
    df['interest_rate_squared'] = df['interest_rate'] ** 2

    # 4. RISK SCORES
    df['risk_score_v1'] = (df['debt_to_income_ratio'] * 0.25 + df['loan_to_income_ratio'] * 0.30 + df['credit_risk_score'] * 0.30 + (df['interest_rate'] / df['annual_income']) * 0.25)
    df['risk_score_v2'] = (df['payment_to_income_ratio'] * 0.40 + df['new_debt_to_income_ratio'] * 0.25 + df['interest_burden'] * 0.25)
```

```

df['affordability_score'] = (df['credit_score_normalized'] * 0.40 +
                             (1 - df['debt_to_income_ratio']) * 0.30 +
                             df['disposable_income_ratio'] * 0.30)
df['financial_health_score'] = df['affordability_score'] * 0.60 - df['risk_score']

# 5. LOAN AMOUNT FEATURES
df['loan_size'] = pd.cut(df['loan_amount'], bins=[0, 10000, 20000, 30000, np.inf],
                         labels=['small', 'medium', 'large', 'very_large'])
df['loan_amount_squared'] = df['loan_amount'] ** 2
df['loan_amount_log'] = np.log1p(df['loan_amount'])
df['annual_income_log'] = np.log1p(df['annual_income'])
df['loan_amount_sqrt'] = np.sqrt(df['loan_amount'])

# 6. BINNING FEATURES
df['income_decile'] = pd.qcut(df['annual_income'], q=10, labels=False, duplicates='drop')
df['credit_decile'] = pd.qcut(df['credit_score'], q=10, labels=False, duplicates='drop')
df['loan_decile'] = pd.qcut(df['loan_amount'], q=10, labels=False, duplicates='drop')
df['dti_decile'] = pd.qcut(df['debt_to_income_ratio'], q=10, labels=False, duplicates='drop')
df['interest_decile'] = pd.qcut(df['interest_rate'], q=10, labels=False, duplicates='drop')

# 7. INTERACTION FEATURES
df['income_x_credit'] = df['annual_income'] * df['credit_score']
df['dti_x_interest'] = df['debt_to_income_ratio'] * df['interest_rate']
df['loan_x_interest'] = df['loan_amount'] * df['interest_rate']
df['income_x_dti'] = df['annual_income'] * df['debt_to_income_ratio']
df['income_credit_loan'] = (df['annual_income'] * df['credit_score']) / (df['loan_amount'])
df['dti_interest_credit'] = (df['debt_to_income_ratio'] * df['interest_rate']) / (df['credit_score'])

# 8. GRADE FEATURES
df['grade'] = df['grade_subgrade'].str[0]
df['subgrade_num'] = pd.to_numeric(df['grade_subgrade'].str[1:], errors='coerce')
grade_map = {'A': 1, 'B': 2, 'C': 3, 'D': 4, 'E': 5, 'F': 6, 'G': 7}
df['grade_numeric'] = df['grade'].map(grade_map)
df['full_grade_score'] = df['grade_numeric'] * 10 + df['subgrade_num']
df['grade_credit_ratio'] = df['full_grade_score'] / (df['credit_score'] / 100)

# 9. STATISTICAL AGGREGATIONS
financial_metrics = ['debt_to_income_ratio', 'loan_to_income_ratio', 'payment_to_income']
df['mean_financial_metrics'] = df[financial_metrics].mean(axis=1)
df['max_financial_burden'] = df[financial_metrics].max(axis=1)
df['min_financial_burden'] = df[financial_metrics].min(axis=1)
df['std_financial_metrics'] = df[financial_metrics].std(axis=1)

# 10. CATEGORICAL COMBINATIONS
df['gender_marital'] = df['gender'] + '_' + df['marital_status']
df['education_employment'] = df['education_level'] + '_' + df['employment_status']
df['gender_education'] = df['gender'] + '_' + df['education_level']
df['marital_employment'] = df['marital_status'] + '_' + df['employment_status']
df['purpose_grade'] = df['loan_purpose'] + '_' + df['grade']
df['employment_purpose'] = df['employment_status'] + '_' + df['loan_purpose']

# 11. ANOMALY FLAGS
df['extreme_dti'] = (df['debt_to_income_ratio'] > df['debt_to_income_ratio'].quantile(0.95))
df['low_income'] = (df['annual_income'] < df['annual_income'].quantile(0.25)).any()
df['large_loan'] = (df['loan_amount'] > df['loan_amount'].quantile(0.75)).astype(bool)
df['risky_combo_1'] = ((df['debt_to_income_ratio'] > 0.4) & (df['credit_score'] < 600))

```

```

df['risky_combo_2'] = ((df['loan_to_income_ratio'] > 0.5) & (df['interest_rate'] < 0.08))
df['safe_combo'] = ((df['credit_score'] > 750) & (df['debt_to_income_ratio'] < 0.5))
df['high_risk_all'] = (df['extreme_dti'] & df['risky_combo_1']).astype(int)

return df

```

```

In [5]: print("--- 3. Applying Feature Engineering & Encoding ---")

# Apply feature engineering
train_fe = complete_feature_engineering(train)
test_fe = complete_feature_engineering(test)

# Encode categorical features
print("\nENCODING CATEGORICAL FEATURES")
categorical_features = train_fe.select_dtypes(include=['object', 'category']).columns
le_dict = {}
for col in categorical_features:
    le = LabelEncoder()
    # Combine train and test for a full fit, ensuring all categories are known
    all_values = pd.concat([train_fe[col].astype(str), test_fe[col].astype(str)])
    le.fit(all_values)

    train_fe[col] = le.transform(train_fe[col].astype(str))
    test_fe[col] = le.transform(test_fe[col].astype(str))
    le_dict[col] = le
    print(f"✓ {col}: {len(le.classes_)} classes")

# Prepare final datasets
print("\nFINAL DATA READY")
feature_cols = [col for col in train_fe.columns if col not in ['id', config.TARGET]]

# Align columns - crucial if FE created different columns
train_cols = set(train_fe.columns)
test_cols = set(test_fe.columns)

missing_in_test = list(train_cols - test_cols - {'id', config.TARGET})
for col in missing_in_test:
    if col in feature_cols:
        test_fe[col] = 0

missing_in_train = list(test_cols - train_cols - {'id'})
for col in missing_in_train:
    if col in feature_cols:
        train_fe[col] = 0

# Ensure final feature list is identical
feature_cols = [col for col in feature_cols if col in test_fe.columns]
X = train_fe[feature_cols]
y = train_fe[config.TARGET]
X_test = test_fe[feature_cols]
test_ids = test_fe['id']

# Fill any remaining NaNs from FE (e.g., from ratios)
X = X.fillna(-1)
X_test = X_test.fillna(-1)

```

```
print(f"X: {X.shape}")
print(f"y: {y.shape}")
print(f"X_test: {X_test.shape}")
print(f"Features: {len(feature_cols)}")
print("\n--- 4. Model Training (LGBM) ---")
```

--- 3. Applying Feature Engineering & Encoding ---

ENCODING CATEGORICAL FEATURES

- ✓ gender: 3 classes
- ✓ marital_status: 4 classes
- ✓ education_level: 5 classes
- ✓ employment_status: 5 classes
- ✓ loan_purpose: 8 classes
- ✓ grade_subgrade: 30 classes
- ✓ credit_category: 5 classes
- ✓ loan_size: 4 classes
- ✓ grade: 6 classes
- ✓ gender_marital: 12 classes
- ✓ education_employment: 25 classes
- ✓ gender_education: 15 classes
- ✓ marital_employment: 20 classes
- ✓ purpose_grade: 48 classes
- ✓ employment_purpose: 40 classes

FINAL DATA READY

```
X: (593994, 84)
y: (593994,)
X_test: (254569, 84)
Features: 84
```

--- 4. Model Training (LGBM) ---

```
In [6]: def train_lightgbm(X, y, X_test, n_splits=5):
    """ Trains LightGBM model using StratifiedKFold """
    skf = StratifiedKFold(n_splits=n_splits, shuffle=True, random_state=config.SEED)
    oof_preds = np.zeros(len(X))
    test_preds = np.zeros(len(X_test))
    feature_importance = pd.DataFrame()

    # Parameters from source notebook
    params = {
        'objective': 'binary',
        'metric': 'auc',
        'boosting_type': 'gbdt',
        'num_leaves': 31,
        'learning_rate': 0.05,
        'feature_fraction': 0.8,
        'bagging_fraction': 0.8,
        'bagging_freq': 5,
        'max_depth': -1,
        'min_child_samples': 20,
        'reg_alpha': 0.1,
        'reg_lambda': 0.1,
        'random_state': config.SEED,
        'verbose': -1,
```

```

        'n_jobs': -1
    }

fold_scores = []

for fold, (train_idx, val_idx) in enumerate(skf.split(X, y)):
    print(f"\n--- Fold {fold + 1}/{n_splits} ---")
    X_train, X_val = X.iloc[train_idx], X.iloc[val_idx]
    y_train, y_val = y.iloc[train_idx], y.iloc[val_idx]

    train_data = lgb.Dataset(X_train, label=y_train)
    val_data = lgb.Dataset(X_val, label=y_val, reference=train_data)

    model = lgb.train(
        params,
        train_data,
        num_boost_round=2000,
        valid_sets=[train_data, val_data],
        callbacks=[lgb.early_stopping(100), lgb.log_evaluation(200)]
    )

    val_preds = model.predict(X_val, num_iteration=model.best_iteration)
    oof_preds[val_idx] = val_preds
    test_preds += model.predict(X_test, num_iteration=model.best_iteration) / n

    score = roc_auc_score(y_val, val_preds)
    fold_scores.append(score)
    print(f"Fold {fold + 1} AUC: {score:.6f}")

    fold_importance_df = pd.DataFrame()
    fold_importance_df["feature"] = X.columns
    fold_importance_df["importance"] = model.feature_importance(importance_type="split")
    fold_importance_df["fold"] = fold + 1
    feature_importance = pd.concat([feature_importance, fold_importance_df], axis=0)

overall_score = roc_auc_score(y, oof_preds)
print(f"\nLightGBM OOF AUC: {overall_score:.6f}")
print(f"Mean: {np.mean(fold_scores):.6f} (+/- {np.std(fold_scores):.6f})")
return oof_preds, test_preds, feature_importance, overall_score

```

In [7]:

```

print("\nTraining LightGBM...")
lgb_oof, lgb_test, lgb_importance, lgb_score = train_lightgbm(X, y, X_test, n_split
print("\n--- 5. Model Training (XGBoost) ---")

```

Training LightGBM...

--- Fold 1/5 ---

Training until validation scores don't improve for 100 rounds

| | |
|---------------------------------|-------------------------|
| [200] training's auc: 0.922442 | valid_1's auc: 0.920568 |
| [400] training's auc: 0.927888 | valid_1's auc: 0.921743 |
| [600] training's auc: 0.932566 | valid_1's auc: 0.922435 |
| [800] training's auc: 0.936837 | valid_1's auc: 0.922586 |
| [1000] training's auc: 0.940291 | valid_1's auc: 0.922635 |

Early stopping, best iteration is:

| | |
|--------------------------------|-------------------------|
| [938] training's auc: 0.939342 | valid_1's auc: 0.922691 |
|--------------------------------|-------------------------|

Fold 1 AUC: 0.922691

--- Fold 2/5 ---

Training until validation scores don't improve for 100 rounds

| | |
|--------------------------------|-------------------------|
| [200] training's auc: 0.922359 | valid_1's auc: 0.91966 |
| [400] training's auc: 0.928167 | valid_1's auc: 0.921101 |
| [600] training's auc: 0.932599 | valid_1's auc: 0.921701 |

Early stopping, best iteration is:

| | |
|--------------------------------|-------------------------|
| [631] training's auc: 0.933185 | valid_1's auc: 0.921781 |
|--------------------------------|-------------------------|

Fold 2 AUC: 0.921781

--- Fold 3/5 ---

Training until validation scores don't improve for 100 rounds

| | |
|--------------------------------|-------------------------|
| [200] training's auc: 0.923076 | valid_1's auc: 0.918779 |
| [400] training's auc: 0.928423 | valid_1's auc: 0.919877 |
| [600] training's auc: 0.932904 | valid_1's auc: 0.920312 |
| [800] training's auc: 0.936921 | valid_1's auc: 0.920498 |

Early stopping, best iteration is:

| | |
|--------------------------------|-------------------------|
| [747] training's auc: 0.935897 | valid_1's auc: 0.920619 |
|--------------------------------|-------------------------|

Fold 3 AUC: 0.920619

--- Fold 4/5 ---

Training until validation scores don't improve for 100 rounds

| | |
|--------------------------------|-------------------------|
| [200] training's auc: 0.922797 | valid_1's auc: 0.919481 |
| [400] training's auc: 0.928375 | valid_1's auc: 0.920635 |
| [600] training's auc: 0.932944 | valid_1's auc: 0.921098 |
| [800] training's auc: 0.936856 | valid_1's auc: 0.921253 |

Early stopping, best iteration is:

| | |
|--------------------------------|-----------------------|
| [835] training's auc: 0.937522 | valid_1's auc: 0.9213 |
|--------------------------------|-----------------------|

Fold 4 AUC: 0.921300

--- Fold 5/5 ---

Training until validation scores don't improve for 100 rounds

| | |
|--------------------------------|-------------------------|
| [200] training's auc: 0.922951 | valid_1's auc: 0.919204 |
| [400] training's auc: 0.928501 | valid_1's auc: 0.920553 |
| [600] training's auc: 0.932888 | valid_1's auc: 0.92081 |

Early stopping, best iteration is:

| | |
|--------------------------------|-------------------------|
| [680] training's auc: 0.934556 | valid_1's auc: 0.920939 |
|--------------------------------|-------------------------|

Fold 5 AUC: 0.920939

LightGBM OOF AUC: 0.921461
Mean: 0.921466 (+/- 0.000724)

--- 5. Model Training (XGBoost) ---

```
In [8]: def train_xgboost(X, y, X_test, n_splits=5):
    """ Trains XGBoost model using StratifiedKFold"""
    skf = StratifiedKFold(n_splits=n_splits, shuffle=True, random_state=config.SEED)
    oof_preds = np.zeros(len(X))
    test_preds = np.zeros(len(X_test))

    # Parameters from source notebook
    params = {
        'objective': 'binary:logistic',
        'eval_metric': 'auc',
        'max_depth': 6,
        'learning_rate': 0.05,
        'subsample': 0.8,
        'colsample_bytree': 0.8,
        'min_child_weight': 1,
        'reg_alpha': 0.1,
        'reg_lambda': 0.1,
        'random_state': config.SEED,
        'tree_method': 'hist',
        'n_jobs': -1
    }

    fold_scores = []

    for fold, (train_idx, val_idx) in enumerate(skf.split(X, y)):
        print(f"\n--- Fold {fold + 1}/{n_splits} ---")
        X_train, X_val = X.iloc[train_idx], X.iloc[val_idx]
        y_train, y_val = y.iloc[train_idx], y.iloc[val_idx]

        model = xgb.XGBClassifier(**params, n_estimators=2000)
        model.fit(
            X_train, y_train,
            eval_set=[(X_val, y_val)],
            early_stopping_rounds=100,
            verbose=200
        )

        val_preds = model.predict_proba(X_val)[:, 1]
        oof_preds[val_idx] = val_preds
        test_preds += model.predict_proba(X_test)[:, 1] / n_splits

        score = roc_auc_score(y_val, val_preds)
        fold_scores.append(score)
        print(f"Fold {fold + 1} AUC: {score:.6f}")

    overall_score = roc_auc_score(y, oof_preds)
    print(f"\nXGBoost OOF AUC: {overall_score:.6f}")
    print(f"Mean: {np.mean(fold_scores):.6f} (+/- {np.std(fold_scores):.6f})")
    return oof_preds, test_preds, overall_score
```

```
In [9]: print("\nTraining XGBoost...")
xgb_oof, xgb_test, xgb_score = train_xgboost(X, y, X_test, n_splits=config.N_SPLITS
print("\n--- 6. Model Training (CatBoost) ---")
```

```
Training XGBoost...
```

```
--- Fold 1/5 ---  
[0] validation_0-auc:0.90833  
[200] validation_0-auc:0.91925  
[400] validation_0-auc:0.92142  
[600] validation_0-auc:0.92198  
[776] validation_0-auc:0.92207  
Fold 1 AUC: 0.922122
```

```
--- Fold 2/5 ---  
[0] validation_0-auc:0.90679  
[200] validation_0-auc:0.91852  
[400] validation_0-auc:0.92065  
[600] validation_0-auc:0.92140  
[772] validation_0-auc:0.92146  
Fold 2 AUC: 0.921505
```

```
--- Fold 3/5 ---  
[0] validation_0-auc:0.90673  
[200] validation_0-auc:0.91692  
[400] validation_0-auc:0.91900  
[600] validation_0-auc:0.91959  
[800] validation_0-auc:0.91970  
[965] validation_0-auc:0.91966  
Fold 3 AUC: 0.919772
```

```
--- Fold 4/5 ---  
[0] validation_0-auc:0.90693  
[200] validation_0-auc:0.91814  
[400] validation_0-auc:0.92006  
[600] validation_0-auc:0.92067  
[800] validation_0-auc:0.92088  
[950] validation_0-auc:0.92083  
Fold 4 AUC: 0.920934
```

```
--- Fold 5/5 ---  
[0] validation_0-auc:0.90723  
[200] validation_0-auc:0.91757  
[400] validation_0-auc:0.91968  
[600] validation_0-auc:0.92023  
[800] validation_0-auc:0.92026  
[919] validation_0-auc:0.92027  
Fold 5 AUC: 0.920333
```

```
XGBoost OOF AUC: 0.920922  
Mean: 0.920933 (+/- 0.000830)
```

```
--- 6. Model Training (CatBoost) ---
```

```
In [10]: def train_catboost(X, y, X_test, n_splits=5):  
    """ Trains CatBoost model using StratifiedKFold """  
    skf = StratifiedKFold(n_splits=n_splits, shuffle=True, random_state=config.SEED)  
    oof_preds = np.zeros(len(X))  
    test_preds = np.zeros(len(X_test))
```

```

# Parameters from source notebook
params = {
    'iterations': 2000,
    'learning_rate': 0.05,
    'depth': 6,
    'l2_leaf_reg': 3,
    'random_seed': config.SEED,
    'loss_function': 'Logloss',
    'eval_metric': 'AUC',
    'early_stopping_rounds': 100,
    'verbose': 200,
    'task_type': 'CPU' # Source notebook uses CPU
}

fold_scores = []

for fold, (train_idx, val_idx) in enumerate(skf.split(X, y)):
    print(f"\n--- Fold {fold + 1}/{n_splits} ---")
    X_train, X_val = X.iloc[train_idx], X.iloc[val_idx]
    y_train, y_val = y.iloc[train_idx], y.iloc[val_idx]

    model = CatBoostClassifier(**params)
    model.fit(X_train, y_train, eval_set=(X_val, y_val), use_best_model=True)

    val_preds = model.predict_proba(X_val)[:, 1]
    oof_preds[val_idx] = val_preds
    test_preds += model.predict_proba(X_test)[:, 1] / n_splits

    score = roc_auc_score(y_val, val_preds)
    fold_scores.append(score)
    print(f"Fold {fold + 1} AUC: {score:.6f}")

overall_score = roc_auc_score(y, oof_preds)
print(f"\nCatBoost OOF AUC: {overall_score:.6f}")
print(f"Mean: {np.mean(fold_scores):.6f} (+/- {np.std(fold_scores):.6f})")
return oof_preds, test_preds, overall_score

```

```

In [11]: print("\nTraining CatBoost...")
cat_oof, cat_test, cat_score = train_catboost(X, y, X_test, n_splits=config.N_SPLIT
print("\n--- 7. Model Evaluation & Ensemble ---")

```

Training CatBoost...

--- Fold 1/5 ---

| | | | |
|-------|----------------------------------------|---------------|-------------------|
| 0: | test: 0.9022510 best: 0.9022510 (0) | total: 178ms | remaining: 5m 56s |
| 200: | test: 0.9166465 best: 0.9166465 (200) | total: 21.1s | remaining: 3m 8s |
| 400: | test: 0.9190007 best: 0.9190007 (400) | total: 41.2s | remaining: 2m 44s |
| 600: | test: 0.9202509 best: 0.9202509 (600) | total: 1m 1s | remaining: 2m 23s |
| 800: | test: 0.9209301 best: 0.9209302 (798) | total: 1m 22s | remaining: 2m 2s |
| 1000: | test: 0.9215918 best: 0.9215918 (1000) | total: 1m 42s | remaining: 1m 42s |
| 1200: | test: 0.9219187 best: 0.9219188 (1198) | total: 2m 3s | remaining: 1m 21s |
| 1400: | test: 0.9221681 best: 0.9221732 (1393) | total: 2m 23s | remaining: 1m 1s |
| 1600: | test: 0.9224090 best: 0.9224090 (1600) | total: 2m 44s | remaining: 40.9s |
| 1800: | test: 0.9225697 best: 0.9225697 (1800) | total: 3m 4s | remaining: 20.4s |
| 1999: | test: 0.9226961 best: 0.9226968 (1997) | total: 3m 25s | remaining: 0us |

bestTest = 0.9226967803

bestIteration = 1997

Shrink model to first 1998 iterations.

Fold 1 AUC: 0.922697

--- Fold 2/5 ---

| | | | |
|-------|----------------------------------------|---------------|-------------------|
| 0: | test: 0.9023162 best: 0.9023162 (0) | total: 108ms | remaining: 3m 35s |
| 200: | test: 0.9160644 best: 0.9160644 (200) | total: 21s | remaining: 3m 8s |
| 400: | test: 0.9184948 best: 0.9184949 (399) | total: 41.4s | remaining: 2m 44s |
| 600: | test: 0.9196654 best: 0.9196654 (600) | total: 1m 1s | remaining: 2m 23s |
| 800: | test: 0.9205471 best: 0.9205471 (800) | total: 1m 21s | remaining: 2m 2s |
| 1000: | test: 0.9209918 best: 0.9209922 (999) | total: 1m 42s | remaining: 1m 42s |
| 1200: | test: 0.9215138 best: 0.9215138 (1200) | total: 2m 2s | remaining: 1m 21s |
| 1400: | test: 0.9218013 best: 0.9218013 (1400) | total: 2m 23s | remaining: 1m 1s |
| 1600: | test: 0.9220076 best: 0.9220076 (1600) | total: 2m 44s | remaining: 41s |
| 1800: | test: 0.9221596 best: 0.9221596 (1800) | total: 3m 4s | remaining: 20.4s |
| 1999: | test: 0.9223127 best: 0.9223127 (1999) | total: 3m 25s | remaining: 0us |

bestTest = 0.922312679

bestIteration = 1999

Fold 2 AUC: 0.922313

--- Fold 3/5 ---

| | | | |
|-------|----------------------------------------|---------------|-------------------|
| 0: | test: 0.9001268 best: 0.9001268 (0) | total: 106ms | remaining: 3m 31s |
| 200: | test: 0.9147979 best: 0.9147979 (200) | total: 21.1s | remaining: 3m 8s |
| 400: | test: 0.9170988 best: 0.9170988 (400) | total: 41.2s | remaining: 2m 44s |
| 600: | test: 0.9182407 best: 0.9182407 (600) | total: 1m 1s | remaining: 2m 24s |
| 800: | test: 0.9190087 best: 0.9190092 (799) | total: 1m 22s | remaining: 2m 3s |
| 1000: | test: 0.9194972 best: 0.9194972 (1000) | total: 1m 43s | remaining: 1m 43s |
| 1200: | test: 0.9199853 best: 0.9199918 (1190) | total: 2m 4s | remaining: 1m 22s |
| 1400: | test: 0.9202397 best: 0.9202403 (1399) | total: 2m 25s | remaining: 1m 2s |
| 1600: | test: 0.9204282 best: 0.9204390 (1591) | total: 2m 46s | remaining: 41.5s |
| 1800: | test: 0.9205645 best: 0.9205698 (1794) | total: 3m 7s | remaining: 20.8s |
| 1999: | test: 0.9206759 best: 0.9206772 (1997) | total: 3m 28s | remaining: 0us |

bestTest = 0.9206772067

bestIteration = 1997

Shrink model to first 1998 iterations.

Fold 3 AUC: 0.920677

--- Fold 4/5 ---

| Iteration | test | best | total | remaining |
|-----------|-----------|------------------|--------|-----------|
| 0 | 0.9006072 | 0.9006072 (0) | 107ms | 3m 33s |
| 200 | 0.9155320 | 0.9155320 (200) | 22.2s | 3m 18s |
| 400 | 0.9180183 | 0.9180183 (400) | 45.6s | 3m 2s |
| 600 | 0.9191428 | 0.9191446 (599) | 1m 7s | 2m 38s |
| 800 | 0.9200291 | 0.9200298 (795) | 1m 30s | 2m 16s |
| 1000 | 0.9204769 | 0.9204769 (1000) | 1m 55s | 1m 55s |
| 1200 | 0.9208387 | 0.9208412 (1193) | 2m 18s | 1m 32s |
| 1400 | 0.9210745 | 0.9210775 (1388) | 2m 40s | 1m 8s |
| 1600 | 0.9212991 | 0.9212991 (1600) | 3m 4s | 45.9s |
| 1800 | 0.9214484 | 0.9214515 (1793) | 3m 27s | 22.9s |
| 1999 | 0.9215817 | 0.9215839 (1960) | 3m 50s | 0us |

bestTest = 0.9215838764

bestIteration = 1960

Shrink model to first 1961 iterations.

Fold 4 AUC: 0.921584

--- Fold 5/5 ---

| Iteration | test | best | total | remaining |
|-----------|-----------|------------------|--------|-----------|
| 0 | 0.9010101 | 0.9010101 (0) | 119ms | 3m 58s |
| 200 | 0.9152413 | 0.9152413 (200) | 22s | 3m 17s |
| 400 | 0.9175739 | 0.9175739 (400) | 43.1s | 2m 52s |
| 600 | 0.9187326 | 0.9187326 (600) | 1m 4s | 2m 30s |
| 800 | 0.9194385 | 0.9194398 (799) | 1m 26s | 2m 9s |
| 1000 | 0.9199161 | 0.9199212 (993) | 1m 49s | 1m 49s |
| 1200 | 0.9202181 | 0.9202221 (1194) | 2m 11s | 1m 27s |
| 1400 | 0.9204521 | 0.9204531 (1398) | 2m 34s | 1m 5s |
| 1600 | 0.9206681 | 0.9206790 (1592) | 2m 55s | 43.6s |
| 1800 | 0.9207619 | 0.9207637 (1798) | 3m 16s | 21.7s |
| 1999 | 0.9208691 | 0.9208691 (1999) | 3m 39s | 0us |

bestTest = 0.9208690996

bestIteration = 1999

Fold 5 AUC: 0.920869

CatBoost OOF AUC: 0.921624

Mean: 0.921628 (+/- 0.000787)

--- 7. Model Evaluation & Ensemble ---

In [12]:

```
# Model Comparison
print("\nMODEL COMPARISON")
comparison = pd.DataFrame({
    'Model': ['LightGBM', 'XGBoost', 'CatBoost'],
    'OOF AUC': [lgb_score, xgb_score, cat_score]
}).sort_values('OOF AUC', ascending=False)
print(comparison)

# Create Ensemble
print("\nCREATING ENSEMBLE")

# 1. Simple Average
```

```

simple_oof = (lgb_oof + xgb_oof + cat_oof) / 3
simple_test = (lgb_test + xgb_test + cat_test) / 3
simple_score = roc_auc_score(y, simple_oof)

# 2. Weighted Average
total_auc = lgb_score + xgb_score + cat_score
w_lgb = lgb_score / total_auc
w_xgb = xgb_score / total_auc
w_cat = cat_score / total_auc
weighted_oof = (lgb_oof * w_lgb) + (xgb_oof * w_xgb) + (cat_oof * w_cat)
weighted_test = (lgb_test * w_lgb) + (xgb_test * w_xgb) + (cat_test * w_cat)
weighted_score = roc_auc_score(y, weighted_oof)

# 3. Rank Average
rank_oof = (rankdata(lgb_oof) + rankdata(xgb_oof) + rankdata(cat_oof)) / (3 * len(l
rank_test = (rankdata(lgb_test) + rankdata(xgb_test) + rankdata(cat_test)) / (3 * l
rank_score = roc_auc_score(y, rank_oof)

# Ensemble Results
ensemble_results = pd.DataFrame({
    'Ensemble': ['Simple Average', 'Weighted Average', 'Rank Average'],
    'OOF AUC': [simple_score, weighted_score, rank_score]
}).sort_values('OOF AUC', ascending=False)

print("\nEnsemble Results:")
print(ensemble_results)
print(f"\nWeights: LGB={w_lgb:.3f}, XGB={w_xgb:.3f}, CAT={w_cat:.3f}")

# Choose best
best_idx = ensemble_results['OOF AUC'].idxmax()
best_name = ensemble_results.loc[best_idx, 'Ensemble']
best_score = ensemble_results.loc[best_idx, 'OOF AUC']

if best_name == 'Simple Average':
    final_preds = simple_test
elif best_name == 'Weighted Average':
    final_preds = weighted_test
else:
    final_preds = rank_test

print(f"\nBest Ensemble: {best_name} (AUC: {best_score:.6f})")
print("\n--- 8. Submission ---")

```

```
MODEL COMPARISON
    Model      OOF AUC
2  CatBoost  0.921624
0  LightGBM  0.921461
1  XGBoost   0.920922
```

CREATING ENSEMBLE

Ensemble Results:

```
    Ensemble      OOF AUC
2      Rank Average  0.921893
1  Weighted Average  0.921865
0  Simple Average  0.921865
```

Weights: LGB=0.333, XGB=0.333, CAT=0.333

Best Ensemble: Rank Average (AUC: 0.921893)

--- 8. Submission ---

```
In [13]: # Create submission
# (Uncommented from source to generate the file as per 'Version' format)
submission = pd.DataFrame({
    'id': test_ids,
    config.TARGET: final_preds
})

submission.to_csv('submission.csv', index=False)
print("SUBMISSION CREATED")
print(f"File: submission.csv")
print(f"Shape: {submission.shape}")
print(f"\nPreview:")
display(submission.head(10))
```

```
SUBMISSION CREATED
File: submission.csv
Shape: (254569, 2)
```

Preview:

| | id | loan_paid_back |
|----------|-----------|-----------------------|
| 0 | 593994 | 0.499309 |
| 1 | 593995 | 0.788466 |
| 2 | 593996 | 0.128066 |
| 3 | 593997 | 0.506112 |
| 4 | 593998 | 0.623759 |
| 5 | 593999 | 0.727296 |
| 6 | 594000 | 0.834262 |
| 7 | 594001 | 0.661948 |
| 8 | 594002 | 0.526185 |
| 9 | 594003 | 0.018947 |